

Session 1 Challenge

Mehmet Alican Noyan

November 14, 2020

```
In [1]: # For data handling and visualization
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

1 Linear Regression example

<https://scikit-learn.org/stable/datasets/index.html#diabetes-dataset>

```
In [2]: # Load the dataset
from sklearn import datasets
diabetes = datasets.load_diabetes()
```

```
In [3]: # Get input, output, and their names
X = diabetes.data
y = diabetes.target
names = diabetes.feature_names
```

```
In [4]: # First 5 data points
df = pd.DataFrame(np.concatenate((X, y[:, np.newaxis])), axis=1, columns=names + ['dis'])
df.head()
```

```
Out[4]:
```


	age	sex	bmi	bp	s1	s2	s3	\
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	
2	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194	-0.032356	
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	

	s4	s5	s6	disease progression
0	-0.002592	0.019908	-0.017646	151.0
1	-0.039493	-0.068330	-0.092204	75.0
2	-0.002592	0.002864	-0.025930	141.0
3	0.034309	0.022692	-0.009362	206.0
4	-0.002592	-0.031991	-0.046641	135.0

```
In [5]: # Only one feature
bmi = X[:, np.newaxis, 2]
```

```
In [6]: # Plot bmi vs. target variable
plt.figure(dpi=200)
plt.scatter(bmi, y)
plt.xlabel('bmi')
plt.ylabel('disease progression')
```

```
Out[6]: Text(0, 0.5, 'disease progression')
```



```
In [7]: # Split the dataset into train and test
```

```
X_train = bmi[:300]
y_train = y[:300]
```

```
X_test = bmi[300:]
y_test = y[300:]
```

```
In [8]: # Sanity check the shapes
```

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```


```
Out[8]: ((300, 1), (300,), (142, 1), (142,))
```

```
In [9]: # Plot train and test
```

```
plt.figure(dpi=200)
plt.scatter(X_train, y_train, label='train')
plt.scatter(X_test, y_test, label='test')
```

```
plt.xlabel('bmi')
plt.ylabel('disease progression')
plt.legend()
```

Out [9]: <matplotlib.legend.Legend at 0x11e1bd9e8>



```
In [10]: # Call the linear regression object
from sklearn import linear_model
model1 = linear_model.LinearRegression()
```

```
In [11]: # LEARN = .fit(input, target)
model1.fit(X_train, y_train)
```

Out [11]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)

```
In [12]: # Get predictions for the test set
preds = model1.predict(X_test)
```

```
In [13]: # Check predictions vs. labels for the first 10 data points
preds[:10], y_test[:10]
```

Out [13]: (array([222.78610979, 127.10097299, 183.88116406, 184.93264908,
113.43166773, 112.38018271, 149.1821584 , 120.79206287,
106.07127259, 152.33661346]),
array([275., 65., 198., 236., 253., 124., 44., 172., 114., 142.]))

```

In [14]: # Quantify the error
         preds[:10] - y_test[:10]

Out[14]: array([ -52.21389021,   62.10097299,  -14.11883594,  -51.06735092,
                -139.56833227,  -11.61981729,  105.1821584 ,  -51.20793713,
                 -7.92872741,   10.33661346])

In [15]: # Use absolute values to avoid positive/negative errors canceling each other
         abs(preds[:10] - y_test[:10])

Out[15]: array([ 52.21389021,   62.10097299,   14.11883594,   51.06735092,
                139.56833227,   11.61981729,  105.1821584 ,   51.20793713,
                 7.92872741,   10.33661346])

In [16]: # Find the mean abs error
         np.mean(abs(preds[:10] - y_test[:10]))

Out[16]: 50.53446360349993

In [17]: # Find the mean abs error, for all test data
         np.mean(abs(preds - y_test))

Out[17]: 50.59683538302396


In [18]: # This and many other metrics are available in sklearn
         from sklearn import metrics
         metrics.mean_absolute_error(y_test, preds)

Out[18]: 50.59683538302396

In [19]: # Plot train/test data with fit
         plt.figure(dpi=200)
         plt.scatter(X_train, y_train, label='train')
         plt.plot(X_test, model1.predict(X_test), label='fit')
         plt.scatter(X_test, y_test, label='test')
         plt.xlabel('bmi')
         plt.ylabel('disease progression')
         plt.legend()

Out[19]: <matplotlib.legend.Legend at 0x11ea96f28>

```




2 Deep learning example

```
In [20]: # Download the MNIST dataset
from keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Using TensorFlow backend.

```
In [21]: # Random digit from the training set.
index = np.random.randint(0, 60000)
plt.imshow(x_train[index])
print('Label: ', y_train[index])
```

Label: 3



```
In [22]: # Build the architecture
from keras.models import Sequential
from keras.layers import Conv2D, Flatten, Activation, Dense

model2 = Sequential()
model2.add(Conv2D(32, (3, 3), padding='same', input_shape=(28, 28, 1)))
model2.add(Activation('relu'))
model2.add(Conv2D(32, (3, 3), strides=2, padding='same'))
model2.add(Activation('relu'))
model2.add(Flatten())
model2.add(Dense(10))
model2.add(Activation('softmax'))
model2.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 32)	320
activation_1 (Activation)	(None, 28, 28, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 32)	9248
activation_2 (Activation)	(None, 14, 14, 32)	0
flatten_1 (Flatten)	(None, 6272)	0

```

-----
dense_1 (Dense)                (None, 10)                62730
-----
activation_3 (Activation)      (None, 10)                0
=====
Total params: 72,298
Trainable params: 72,298
Non-trainable params: 0
-----

```

```

In [23]: # Set the optimizer and the loss
         from keras.optimizers import Adam, SGD
         opt = SGD(lr=0.001)
         model2.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

In [24]: # Put data into suitable shape
         from keras.utils import to_categorical
         y_test = to_categorical(y_test)
         y_train = to_categorical(y_train)

         x_train = x_train.reshape(-1, 28, 28, 1)
         x_test = x_test.reshape(-1, 28, 28, 1)


In [25]: # Train and test
         H = model2.fit(x_train, y_train, batch_size=32, epochs=1, validation_data=(x_test, y_test))

Train on 60000 samples, validate on 10000 samples
Epoch 1/1
60000/60000 [=====] - 53s 889us/step - loss: 0.1489 - acc: 0.9710 - val_loss: 0.1489 - val_acc: 0.9710

In [26]: # Random image from the test set
         index = np.random.randint(0, 10000)
         random_image = x_test[np.newaxis, index]
         plt.imshow(random_image.reshape(28, 28))
         print('Label:', np.argmax(y_test[index]))
         print('Prediction: ', np.argmax(model2.predict(random_image))) # model.predict(image)

Label: 8
Prediction: 3

```



3 Challenge 1

```
In [27]: from sklearn import datasets
         diabetes = datasets.load_diabetes()
```

- Repeat the example: predict disease progression from bmi
- Check errors with different metrics e.g. mean squared error (mse) <https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics>
- Compare mse with mae
- Find the intercept and slope using `sklearn.linear_model.LinearRegression` attributes https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html
- Instead of bmi only, use more features and fit again, check if the performance improves
- Use other algorithms e.g. Random Forests

4 Data = Signal + Noise

- Data generation
 - Signal
 - Noise
 - Signal + Noise


```
In [28]: # Newton's second law as signal
```

```
m = 2 # kg
F = 10 # N

a = F/m # m/s2

# a = intercept + slope * F

a
```

```
Out[28]: 5.0
```

```
In [29]: # 10 random F's between 0 and 1
```

```
F = np.random.rand(10, 1)
F
```

```
Out[29]: array([[0.27265566],
               [0.45639664],
               [0.59645395],
               [0.93229163],
               [0.42039381],
               [0.79091579],
               [0.92921021],
               [0.89619034],
               [0.93065236],
               [0.95336649]])
```

```
In [30]: # 10 random F's between 0 and 100
```

```
F = np.random.rand(10, 1)*100
F
```

```
Out[30]: array([[17.48382141],
               [76.99769087],
               [ 9.11378162],
               [89.1453672 ],
               [69.00352034],
               [18.27328023],
               [93.63434462],
               [41.13464976],
               [56.54294344],
               [20.48335529]])
```

```
In [31]: # Calculate a using Newton's second law
```


```
a = F/m
a
```

```
Out[31]: array([[ 8.7419107 ],
               [38.49884543],
```

```
[ 4.55689081],
[44.5726836 ],
[34.50176017],
[ 9.13664012],
[46.81717231],
[20.56732488],
[28.27147172],
[10.24167764]])
```

```
In [32]: # Plot F vs. a
plt.figure(dpi=200)
plt.scatter(F, a)
plt.xlabel('F (N)')
plt.ylabel('a (m/s2)')
```

```
Out[32]: Text(0, 0.5, 'a (m/s2)')
```




```
In [33]: # Let's add noise to the signal
m = 2 # kg
F = np.random.rand(10, 1) * 100 # N
noise = np.random.rand(10, 1) * 10 - 5
a = F/m + noise # m/s2

# Plot F vs. a
plt.figure(dpi=200)
```

```
plt.scatter(F, a)
plt.xlabel('F (N)')
plt.ylabel('a (m/s2)')
```

Out [33]: Text(0, 0.5, 'a (m/s²)')

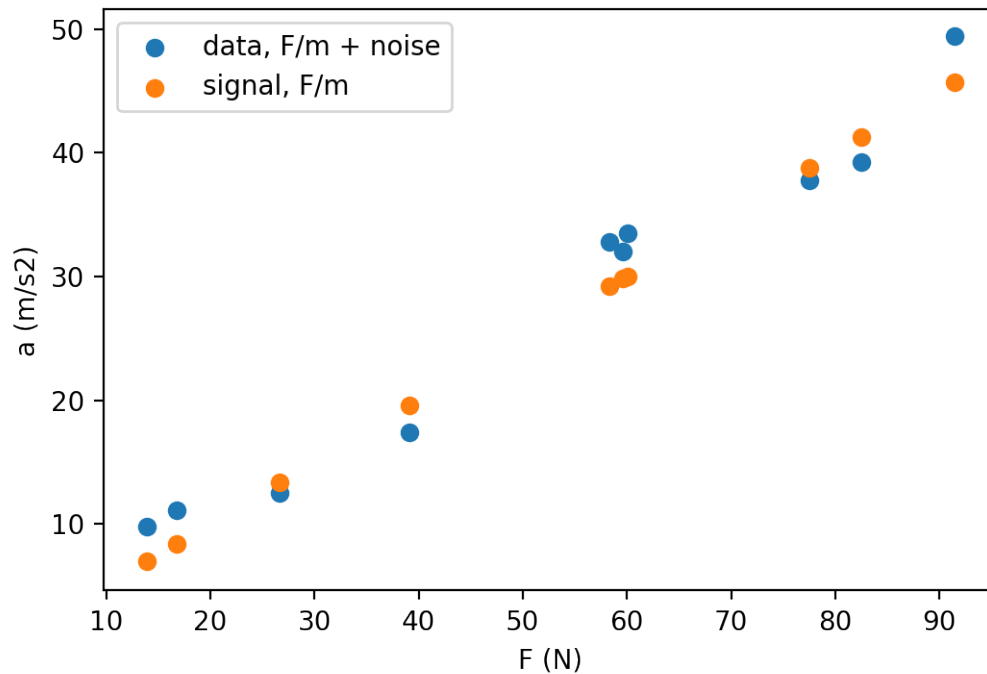


Data = Signal + Noise

$a = F/m + \text{noise}$

```
In [34]: # Signal vs. Data
plt.figure(dpi=200)
plt.scatter(F, a, label='data, F/m + noise')
plt.scatter(F, F/m, label='signal, F/m')
plt.xlabel('F (N)')
plt.ylabel('a (m/s2)')
plt.legend()
```

Out [34]: <matplotlib.legend.Legend at 0x135919eb8>

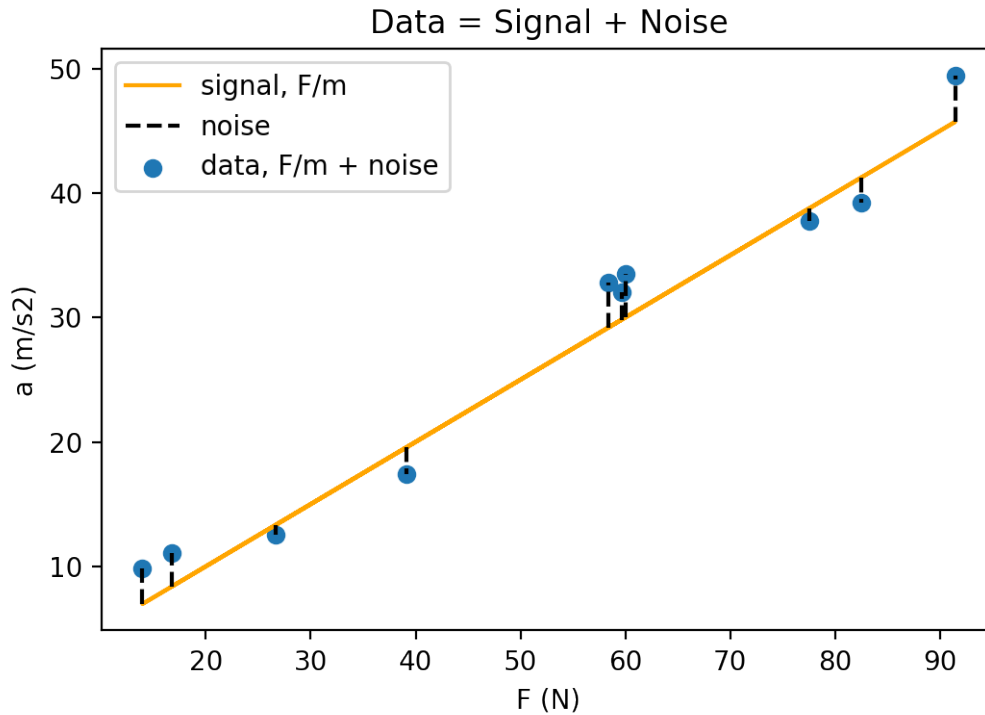


```
In [35]: # Data = Signal + Noise
plt.figure(dpi=200)
plt.scatter(F, a, label='data, F/m + noise')
plt.plot(F, F/m, label='signal, F/m', color='orange')

for i in range(len(F)):
    plt.plot([F[i], F[i]], [F[i]/m, a[i]], 'k--', label='noise' if i == 0 else None)

plt.xlabel('F (N)')
plt.ylabel('a (m/s²)')
plt.legend()
plt.title('Data = Signal + Noise')

Out[35]: Text(0.5, 1.0, 'Data = Signal + Noise')
```



The aim of machine learning is to find the signal

```
In [36]: # We know the signal, let's see if ML can find it
from sklearn.linear_model import LinearRegression
linear_model = LinearRegression()
linear_model.fit(F, a)
```


```
Out[36]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
In [37]: # Signal vs. Fit
plt.figure(dpi=200)
plt.scatter(F, a, label='data')
plt.plot(F, F/m, label='signal', color='orange')
plt.plot(F, linear_model.predict(F), label='fit')

for i in range(len(F)):
    plt.plot([F[i], F[i]], [F[i]/m, a[i]], 'k--', label='noise' if i == 0 else None)

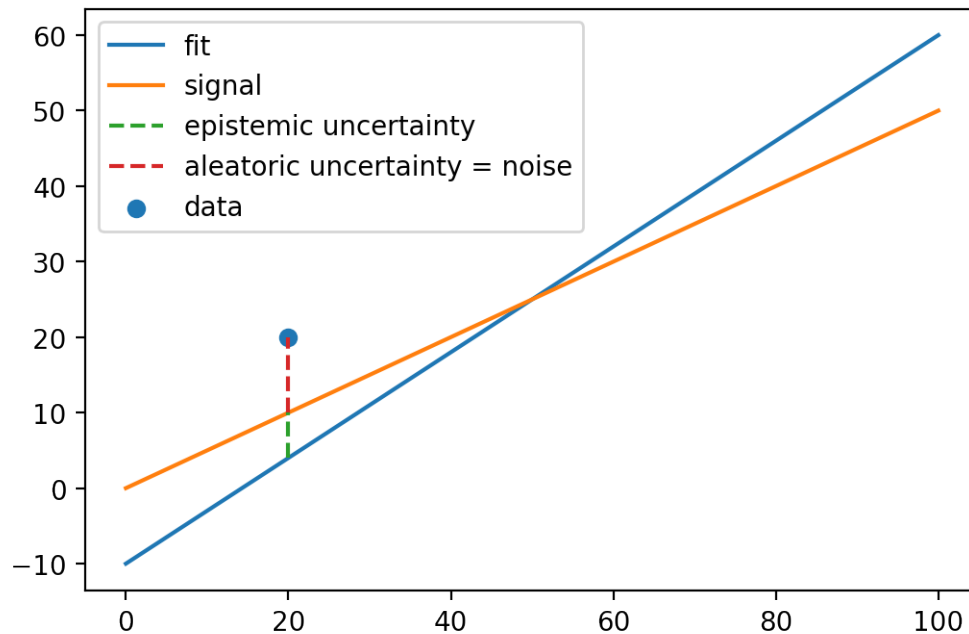
plt.xlabel('F (N)')
plt.ylabel('a (m/s2)')
plt.legend()
```

```
Out[37]: <matplotlib.legend.Legend at 0x136332320>
```



```
In [38]: # Aleatoric vs. Epistemic uncertainty
plt.figure(dpi=200)
plt.plot([0, 100], [-10, 60], label='fit')
plt.plot([0, 100], [0, 50], label='signal')
plt.scatter([20], [20], label='data')
plt.plot([20, 20], [4, 10], '--', label='epistemic uncertainty')
plt.plot([20, 20], [10, 20], '--', label='aleatoric uncertainty = noise')
plt.legend()
```

Out[38]: <matplotlib.legend.Legend at 0x136370198>



5 Challenge 2


- Plot how epistemic and aleatoric uncertainty will change wrt the number of data. (fixed noise)
- Plot how epistemic and aleatoric uncertainty will change wrt the noise. (fixed number of data)
- Make an educated guess before plotting.

6 Overfitting

```
In [39]: # Plot again the prev. data
plt.figure(dpi=200)
plt.scatter(F, a, label='data')
plt.plot(F, F/m, label='signal', color='orange')
plt.plot(F, linear_model.predict(F), label='fit')

plt.xlabel('F (N)')
plt.ylabel('a (m/s²)')
plt.legend()
```

Out[39]: <matplotlib.legend.Legend at 0x10c472a90>



```
In [40]: # Fit a degree 5 polynomial to F and a
z = np.poly1d(np.polyfit(F.reshape(-1), a.reshape(-1), 5))
```


```
In [41]: # Plot both fits
```

```
plt.figure(dpi=200)
plt.scatter(F, a, label='data')
plt.plot(F, F/m, label='signal', color='orange')
plt.plot(F, linear_model.predict(F), label='degree 1')
```

```
plt.plot(np.arange(min(F), max(F)), z(np.arange(min(F), max(F))), label='degree 5', c
```

```
plt.xlabel('F (N)')
plt.ylabel('a (m/s²)')
plt.legend()
```

```
Out[41]: <matplotlib.legend.Legend at 0x136353ba8>
```

In [42]: *# Generate new data*

```
m = 2 # kg
F_new = np.random.rand(10, 1) * (max(F)-min(F)) + min(F) # N
noise = np.random.rand(10, 1) * 10 - 5
a_new = F_new/m + noise # m/s2
```

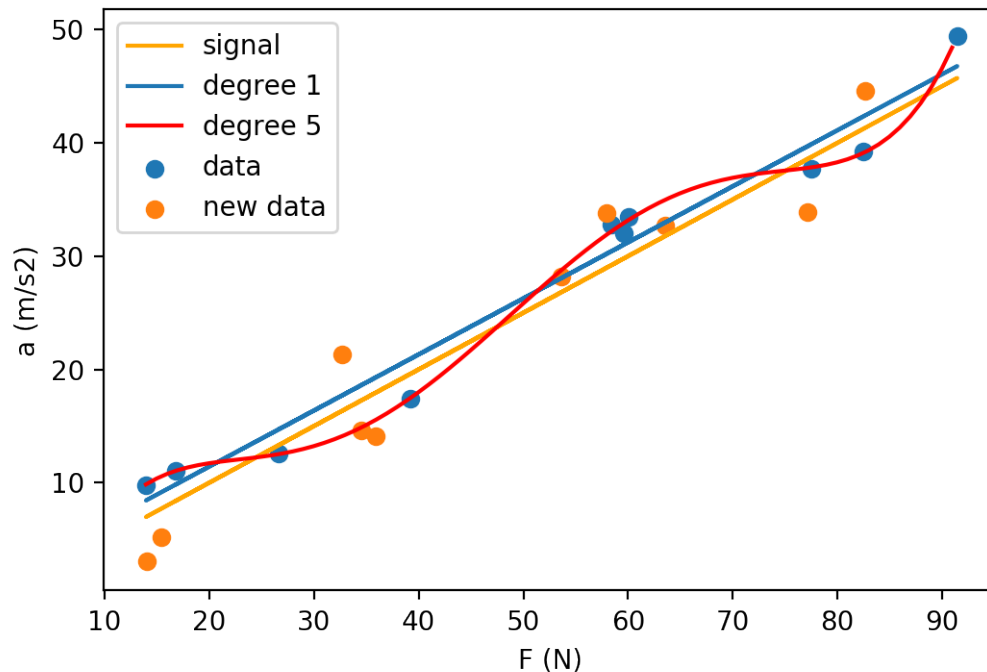
In [43]: *# Both fits on train and test data*

```
plt.figure(dpi=200)
plt.scatter(F, a, label='data')
plt.scatter(F_new, a_new, label='new data')
plt.plot(F, F/m, label='signal', color='orange')
plt.plot(F, linear_model.predict(F), label='degree 1')

plt.plot(np.arange(min(F), max(F)), z(np.arange(min(F), max(F))), label='degree 5', c=

plt.xlabel('F (N)')
plt.ylabel('a (m/s2)')
plt.legend()
```

Out[43]: <matplotlib.legend.Legend at 0x136cd4710>



7 Challenge 3

- Calculate training error for both fits
- Calculate test error for both fits
- Fit different degree polynomials and plot training and test errors wrt degree of the polynomial fit
- Generate data using a quartic signal, repeat previous step, and discover yourself what “underfitting” is.

8 Challenge 4

Implement linear regression using gradient descent

```
In [44]: import numpy as np
import matplotlib.pyplot as plt

In [45]: # Generate data
# 1. Define m and n
m = 10 # number of examples
n = 1 # number of features

In [46]: # 2. Generate X
X = np.random.uniform(size=(m, n))
X
```

```
Out [46]: array([[0.64985039],
                [0.62500976],
                [0.69915575],
                [0.33744647],
                [0.31746079],
                [0.50201737],
                [0.77397666],
                [0.22890078],
                [0.193001  ],
                [0.29300021]])
```

```
In [47]: # 3. Add 1's to X for the bias term for the dot product
X = np.concatenate((np.ones((m, 1)), X), axis=1)
X
```

```
Out [47]: array([[1.          , 0.64985039],
                [1.          , 0.62500976],
                [1.          , 0.69915575],
                [1.          , 0.33744647],
                [1.          , 0.31746079],
                [1.          , 0.50201737],
                [1.          , 0.77397666],
                [1.          , 0.22890078],
                [1.          , 0.193001  ],
                [1.          , 0.29300021]])
```

```
In [48]: # 4. Randomly generate parameters
theta_true = np.random.uniform(size=(n+1, 1))
theta_true
```


```
Out [48]: array([[0.04970101],
                [0.55483883]])
```

```
In [49]: # 5. Calculate y
y = np.dot(X, theta_true)
y
```


```
Out [49]: array([[0.41026324],
                [0.39648069],
                [0.43761977],
                [0.23692942],
                [0.22584059],
                [0.32823974],
                [0.47913332],
                [0.17670406],
                [0.15678546],
                [0.21226891]])
```

```
In [50]: # 6. Plot the generated X and y
plt.figure(dpi=200)
plt.scatter(X[:, 1], y)
```

Out [50]: <matplotlib.collections.PathCollection at 0x1376178d0>



```
In [51]: def plotter(X, y, theta):  
        '''  
        Arguments:  
            n = 1  
            X: np array with shape (m, 1+n)  
            y: np array with shape (m, 1)  
            theta: np array with shape (1+n, 1)  
        Returns:  
            None  
            Plots X, y and the fit defined by theta  
        '''  
        plt.figure(dpi=200)  
        plt.scatter(X[:, 1], y)  
        linex = np.linspace(X[:, 1].min(), X[:, 1].max())  
        liney = theta[0] + theta[1]*linex  
        plt.plot(linex, liney)  
  
In [52]: # 7. Plot the generated X and y, together with the signal used to generate them  
        plotter(X, y, theta_true)
```




8.1 Gradient Descent

```
In [53]: # Random weight initialization
theta = np.random.uniform(size=(n+1, 1))
theta
```

```
Out[53]: array([[0.29412786],
               [0.70486621]])
```

```
In [54]: # X, y together with random fit
plotter(X, y, theta)
```



```
In [55]: # Sub-challenge 1
def cost_function(X, y, theta):
    '''
    Arguments:
        X: np array with shape (m, 1+n)
        y: np array with shape (m, 1)
        theta: np array with shape (1+n, 1)
    Returns:
        cost, a float.
    '''

    return cost

In [56]: # It should be a positive number
cost_function(X, y, theta)

Out[56]: 4.9675336216256785

In [57]: # It should be zero
cost_function(X, y, theta_true)

Out[57]: 0.0

In [58]: # Sub-challenge 2
# Take the theta, update it using the GD rule
def GD_one_step(X, y, theta, lr):
```

```

'''
Updates theta using one gradient descent step.
Arguments:
    X: np array with shape (m, 1+n)
    y: np array with shape (m, 1)
    theta: np array with shape (1+n, 1)
    lr: learning rate, a float
Returns:
    theta_new: np array with shape (1+n, 1), updated after one GD step
'''

return theta_new

In [59]: print('theta: ', theta.T) # Theta before update
        print('theta_true: ', theta_true.T) # True theta

theta: [[0.29412786 0.70486621]]
theta_true: [[0.04970101 0.55483883]]

In [60]: theta = GD_one_step(X, y, theta, lr=0.01)

In [61]: print('theta: ', theta.T) # Theta after the first update
        print('theta_true: ', theta_true.T) # True theta

theta: [[0.2909905 0.70335549]]
theta_true: [[0.04970101 0.55483883]]


In [62]: # Sub-challenge 3
        # Repeat update
def GD(X, y, lr, epoch):
    '''
    Finds theta from X and y using gradient descent.
    Starts by random init. of theta
    Repeats GD_one_step for predefined number of epochs
    Arguments:
        X: np array with shape (m, 1+n)
        y: np array with shape (m, 1)
        lr: learning rate, a float
        epoch: an integer
    Returns:
        theta: np array with shape (1+n, 1), final theta after training
    '''

    return theta

In [63]: # Linear regression with Gradient Descent
        theta = GD(X, y, lr=0.01, epoch=10000)


```

```
In [64]: # Plot the learned theta
plotter(X, y, theta)
```



```
In [65]: def plotter_multiple(X, y, theta, theta_true):
        """
        For comparing the learned theta with the true theta
        Arguments:
            n = 1
            X: np array with shape (m, 1+n)
            y: np array with shape (m, 1)
            theta: np array with shape (1+n, 1)
            theta_true: np array with shape (1+n, 1)
        Returns:
            None
            Plots X, y, fit(theta), signal(theta_true)
        """
        plt.figure(dpi=200)
        plt.scatter(X[:, 1], y)
        line_x = np.array([X[:, 1].min(), X[:, 1].max()])
        plt.plot(line_x, theta[0] + theta[1]*line_x, '--', label='fit')
        plt.plot(line_x, theta_true[0] + theta_true[1]*line_x, ':', label='signal')
        plt.legend()

In [66]: # Fit (learned theta) vs. Signal (true theta)
plotter_multiple(X, y, theta, theta_true)
```

```
In [67]: # Sub-challenge 3 cont.
# Save loss at each epoch for plotting
def GD_memory(X, y, lr, epoch):
    """
    Finds theta from X and y using gradient descent.
    Starts by random init. of theta
    Repeats GD_one_step for predefined number of epochs
    At each epoch calculates the loss and saves it into a list
    Arguments:
        X: np array with shape (m, 1+n)
        y: np array with shape (m, 1)
        lr: learning rate, a float
        epoch: an integer
    Returns:
        theta: np array with shape (1+n, 1), final theta after training
        memory: a list of floats, each item is loss at each epoch
    """
    memory = []


    return theta, memory

In [68]: theta, memory = GD_memory(X, y, lr=0.01, epoch=1000)

In [69]: def loss_plotter(memory):
    """
    Plots the list of losses
```

```
'''  
plt.figure(dpi=200)  
plt.plot(memory)  
plt.ylabel('loss')  
plt.xlabel('number of epochs')
```


```
In [70]: loss_plotter(memory)
```



8.2 It should work with noisy data as well


```
In [71]: # let's introduce noise  
y_real = y * np.random.uniform(0.97, 1.03, m).reshape(-1, 1)
```

```
In [72]: # Plot noisy X and y, with signal  
plotter(X, y_real, theta_true)
```




```
In [73]: theta, memory = GD_memory(X, y_real, lr=0.01, epoch=10000)
```

```
In [74]: loss_plotter(memory)
```



```
In [75]: # Fit vs. Signal for noisy data
plotter_multiple(X, y_real, theta, theta_true)
```



8.3 How about more features and noise

```
In [76]: # Repeat data generation with more features + noise
m = 10 # number of examples
n = 2 # number of features
```

```
In [77]: X = np.random.uniform(size=(m, n))
X
```

```
Out[77]: array([[0.20011324, 0.06261688],
 [0.953954 , 0.24166609],
 [0.07787241, 0.92590823],
 [0.96146176, 0.19724506],
 [0.17035948, 0.99479614],
 [0.00209128, 0.02980649],
 [0.05672694, 0.07903966],
 [0.87311992, 0.01305432],
 [0.12750353, 0.7170062 ],
 [0.41163831, 0.59416031]])
```

```
In [78]: X = np.concatenate((np.ones((m, 1)), X), axis=1)
X
```

```
Out [78]: array([[1.          , 0.20011324, 0.06261688],
                 [1.          , 0.953954   , 0.24166609],
                 [1.          , 0.07787241, 0.92590823],
                 [1.          , 0.96146176, 0.19724506],
                 [1.          , 0.17035948, 0.99479614],
                 [1.          , 0.00209128, 0.02980649],
                 [1.          , 0.05672694, 0.07903966],
                 [1.          , 0.87311992, 0.01305432],
                 [1.          , 0.12750353, 0.7170062 ],
                 [1.          , 0.41163831, 0.59416031]])
```


```
In [79]: theta_true = np.random.uniform(size=(n+1, 1))
        theta_true
```

```
Out [79]: array([[0.10835929],
                 [0.88547444],
                 [0.82740538]])
```

```
In [80]: y = np.dot(X, theta_true)
```

```
In [81]: # Add noise
        y = y * np.random.uniform(0.97, 1.03, m).reshape(-1, 1)
```

```
In [82]: theta, memory = GD_memory(X, y, lr=0.1, epoch=10000)
        loss_plotter(memory)
```



```
In [83]: # Learned params vs. True params
         print('theta: ', theta.T)
         print('theta_true: ', theta_true.T)

theta: [[0.11139771 0.86932284 0.82249641]]
theta_true: [[0.10835929 0.88547444 0.82740538]]

In [84]: # Prediction vs. Label
         print('Prediction:', np.dot(X, theta)[0])
         print('Ground truth:', y[0])

Prediction: [0.33686288]
Ground truth: [0.33565044]

In [ ]:
```